



YONSEI
UNIVERSITY

DyTox: Transformers for Continual Learning with DYnamic TOnken eXpansion [CVPR 2022]


- Young Jo Choi
- Department of Digital Analytics

Severance

Abstract

- A recent trend in continual learning indicates that **dynamic architectures** based on an expansion of the parameters can reduce catastrophic forgetting efficiently. (ex) DER)
- However, they require complex tuning to balance the growing number of parameters and suffer from overhead as they share little information between each task's architecture.
- In this paper, the authors propose a **transformer architecture** based on a dedicated encoder/decoder framework.
(the encoder and decoder are shared among all tasks)

Continual Learning

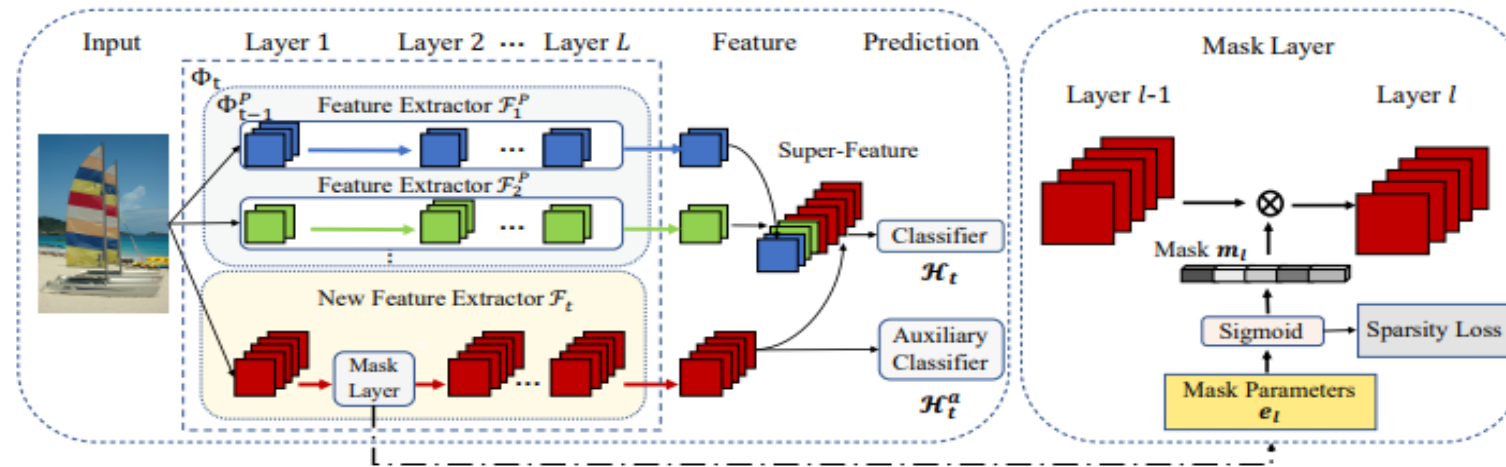


- With new classes of data constantly appearing, simple fine-tuning suffers from **catastrophic forgetting** data which is already learned.
- Continual learning models aim at balancing a **stability/plasticity trade-off** where old data are not forgotten (stability to changes) while learning new incoming data (plasticity to adapt).

Continual Dynamic Networks


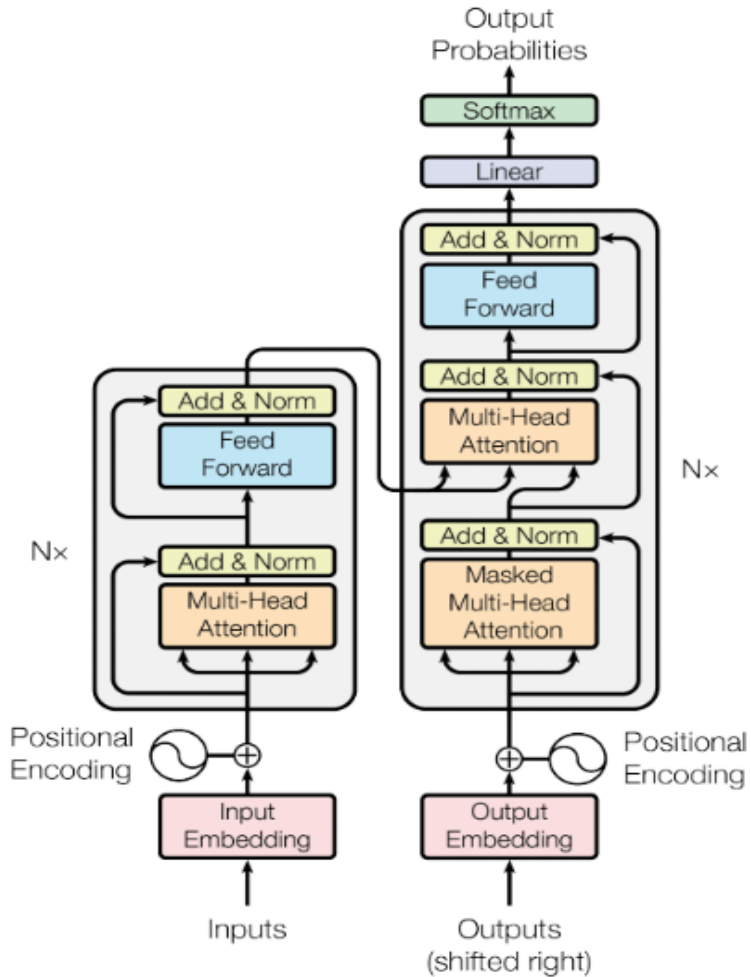
- Be able to cope with growing training distributions by expanding subnetworks for each task.
- However, in this extending method, as the tasks are added, the model size continues to increase linearly. (**memory-overhead**)
- More recently, **DER** has achieved state-of-the-art performance by concatenating the representation of an image by extending a feature extractor per task and using unified classifier to eliminate the need to use task identifiers in testing.

Continual Dynamic Networks



- DER uses a channel-level mask to prune to reduce the growing model size, but it is sensitive to hyperparameters, so the hyperparameters are tuned differently for each experiment.
- Ex) Different hyperparameters are used depending on whether the same dataset is learned by 10 steps or 50 steps.
- Therefore, it can be said to be **unrealistic** in true continuous learning where the number of classes is unknown in advance.

Transformers



1) Get the Q, K, and V vectors from the product of the embedding vector and the weight matrix for each input word.

2) Attention score

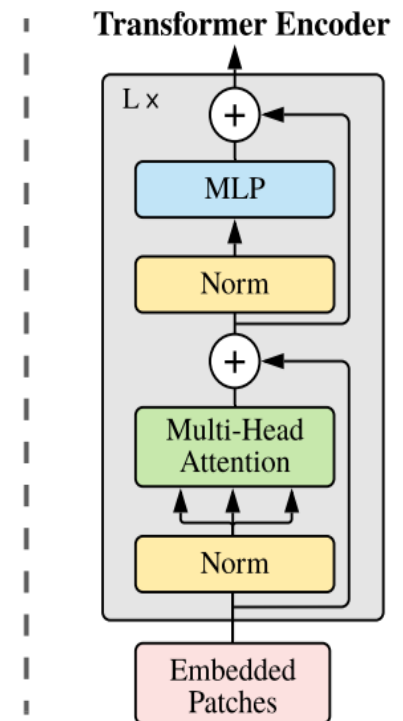
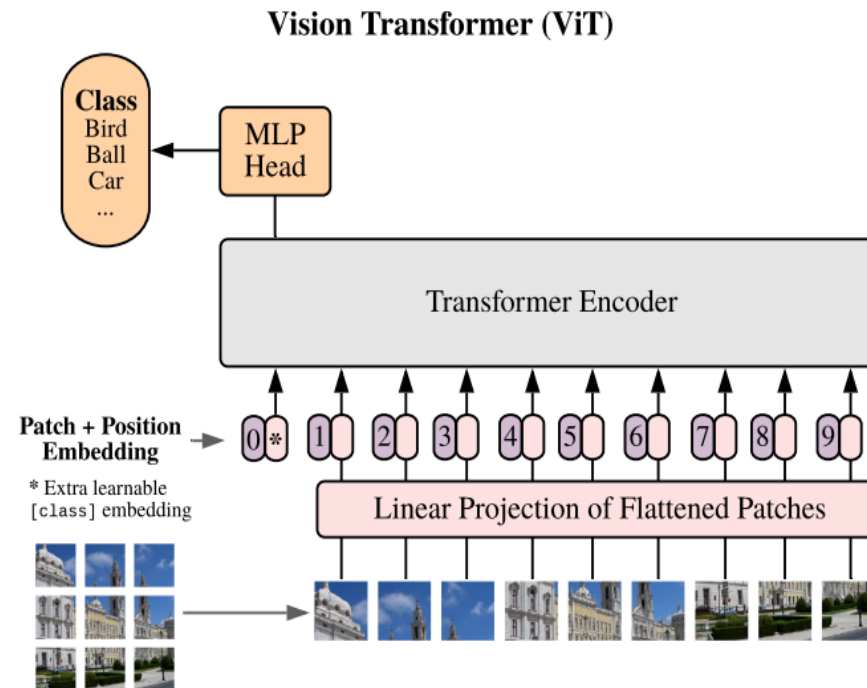
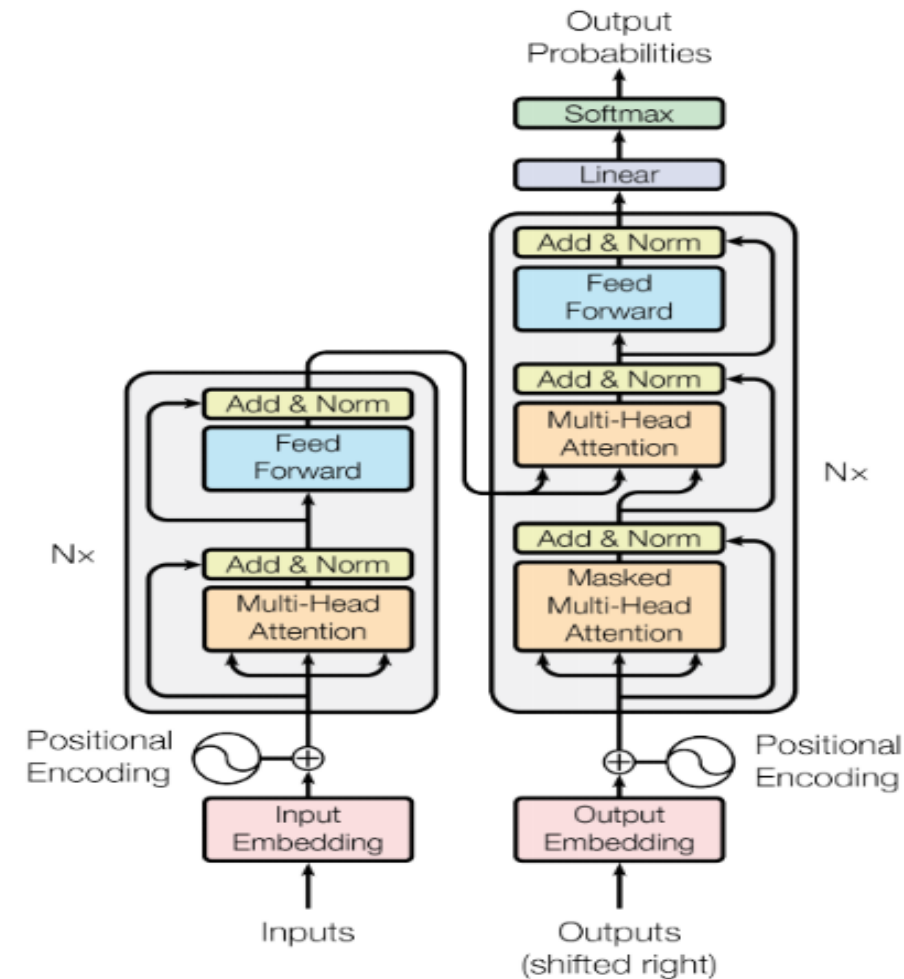
$$Attention(Q, K, V) = Softmax\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

3) Repeat the process with multiple self-attention (head)
: concat and linear transform

4) \rightarrow add & norm \rightarrow MLP
 \rightarrow add & norm

Transformers (ViT)

- Use only the encoder of the transformer
- Split the image into multiple patches (considered as one token)
- Flatten → Linear transformation → results in embedding vectors
- Slightly different order of application in the encoding blocks



DyTox Transformer Model

Setting

- At a given step $t \in \{1 \dots T\}$,
- Training data : $\{(x_i^t, y_i^t)\}_i$
- All task label sets are exclusive : $C^0 \cap C^1 \dots C^T = \emptyset$
- Only a few samples from previous tasks $\{1 \dots t - 1\}$ are available for training step t as **rehearsal** data.
- **Goal** of task t : classify test data coming from all seen classes $C^{1:t}$

Symbol	Meaning
(x_i^t, y_i^t)	Input sample & its label from the t^{th} task
C^t	Label set of the t^{th} task
$C^{1:t}$	All labels from all seen tasks
θ_t	Task token of the t^{th} task
Clf_t	Independent classifier of the t^{th} task
SAB_l	l^{th} Self-Attention Block
TAB	Task-Attention Block

DyTox Transformer Model




Figure 2: **DyTox transformer model.** An image is first split into multiple patches, embedded with a linear projection. The resulting patch tokens are processed by 5 successive Self-Attention Blocks (SAB) (Sec. 3.1). For each task ($t = 1 \dots T$), the processed patch tokens are then given to the Task-Attention Block (TAB) (Sec. 3.2): each forward through the TAB is modified by a different task-specialized token θ_t for $t \in \{1 \dots T\}$ (Sec. 3.3). The T final embeddings are finally given separately to independent classifiers Clf_t each predicting their task's classes C^t . All $|C^{1:T}|$ logits are activated with a sigmoid. For example, at task $t = 3$, one forward is done through the SABs and three task-specific forwards through the unique TAB.

Encoder

- Patch tokenizer
- Self-Attention (SA) block

Decoder

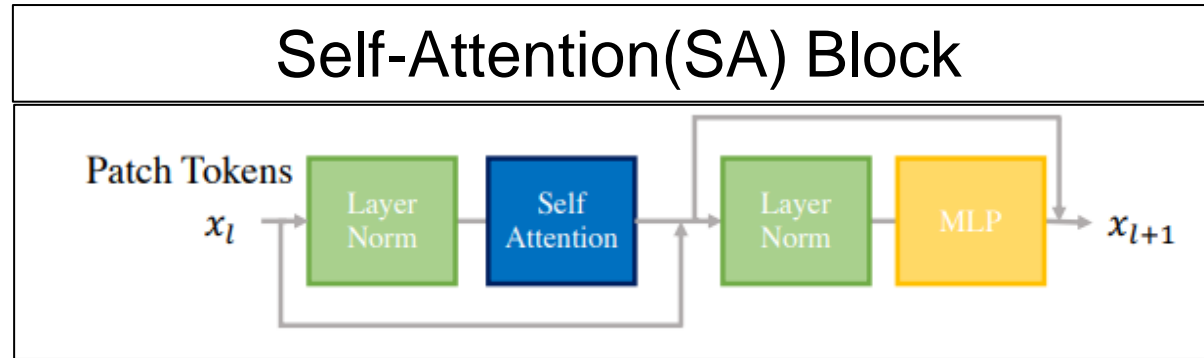
- Task-Attention Block (TAB)
 - With task tokens
- classifier

Structure

- Patch tokenizer
 - fixed-size input RGB image
 - cropped into N patches
 - projected with a linear layer to a dimension D
 - (these operations are done with a single 2D convolution)
 - Result tensor : $x_0 \in \mathbb{R}^{N \times D}$
 - + element-wise sum of positional embedding $p \in \mathbb{R}^{N \times D}$

(In original ViT, classification token is concatenated to the result tensor
→ $[x_{cls}, x_0] \in \mathbb{R}^{(N+1) \times D}$)

Structure



- Self-Attention (SA) based encoder
 - $x'_l = x_l + \text{SA}_l(\text{Norm}_{l,1}(x_l))$
 - $x_{l+1} = x'_l + \text{MLP}_l(\text{Norm}_{l,2}(x'_l))$

Repeat from $l = 1$ to $l = L \rightarrow$ result tensor : $x_L \in \mathbb{R}^{N \times D}$

Structure

- Task-Attention Block
 - In original ViT, “class token” is appended to the patch tokens after tokenizer and is given to a linear classifier with a softmax activation to predict final probabilities.
 - In this paper, “**task token**” (i^{th} task is denoted θ_i) is concatenated to the result of last SA block

$$z_i = [\theta_i, x_L] \in \mathbb{R}^{(N+1) \times D}$$

to exploit this task token, “Task-Attention” layer is newly defined.

Structure

- Task-Attention Block
 - Contrary to the classical Self-Attention, the Task-Attention defines its query (Q_i) only from the task-token θ_i without using the patch tokens x_L

Task-Attention
$Q_i = W_q \theta_i$ $K_i = W_k z_i$ $V_i = W_v z_i$ $A_i = \text{Softmax}(Q_i \cdot K_i^T / \sqrt{d/h})$ $O_i = W_o A_i V_i + b_o \in \mathbb{R}^{1 \times D}$ (output)
d is the embedding dimension, h is the number of attention heads.



Task-Attention Block
$c' = c + \text{TA}(\text{Norm}_1(z))$ $c'' = c' + \text{MLP}(\text{Norm}_2(c'))$

Summarization of all structure
$e_i = \text{TAB} \circ ([\theta_i, \text{SAB}_{l=L} \circ \dots \text{SAB}_{l=1}(x_0)]) \in \mathbb{R}^D$ $\tilde{y}_i = \text{Clf}(e_i) = W_c \text{Norm}_c(e_i) + b_c$

Structure

- Dynamic task token expansion
 - For given image x , through tokenizing and the multiple SABs, x_L is the output of Encoder.
 - In Decoder, each TAB forward passes is executed with a different task token $\theta_i \rightarrow$ different task-specific forward \rightarrow produce task-specific embedding : e_i

$$e_1 = \text{TAB}([\boldsymbol{\theta}_1, x_L])$$

$$e_2 = \text{TAB}([\boldsymbol{\theta}_2, x_L])$$

...

$$e_t = \text{TAB}([\boldsymbol{\theta}_t, x_L])$$

- Rather than concatenating all embeddings $\{e_1, e_2, \dots, e_t\}$, and make one classifier, “**task-specific classifiers**” are suggested.

Structure

- **task-specific classifiers** : $\hat{y}_i = \text{Clf}_i(e_i) = \sigma(W_i \text{Norm}_i e_i + b_i)$
- The predictions for the classes $\hat{y}_i \in \mathcal{C}^i$, where σ is the sigmoid and loss is binary cross-entropy.
- In comparison with the softmax activation, the element-wise sigmoid activation reduces the overconfidence in recent classes \rightarrow the model is better calibrated.

Structure

- The overall structure of the DyTox strategy

Algorithm 1 DyTox's forward pass at step t

Input: x_0 (initial patch tokens), y (ground-truth labels)

Output: $\hat{y}_{1:t}$ (predictions for all classes of $\mathcal{C}^{1:t}$)

```
1:  $x_L \leftarrow \text{SAB}_{l=L} \circ \dots \circ \text{SAB}_{l=1}(x_0)$ 
2: for  $i \leftarrow 1$ ;  $i \leq t$ ;  $i++$  do
3:    $e_i \leftarrow \text{TAB}([\theta_i, x_L])$ 
4:    $\hat{y}_i \leftarrow \text{Clf}_i(e_i)$ 
5: end for
6:  $\hat{y}_{1:t} \leftarrow [\hat{y}_1, \dots, \hat{y}_t]$ 
```

More efficient than a naïve parameter expansion network

- Memory overhead is almost null.
The model's expansion is limited to a new task token per new task (in this case, only $d=384$ new parameters)
← very small compared to the total model size
($\approx 11\text{M}$ parameters)

Structure

- Context
 - The Encoder is shared (both in memory and execution) for all outputs.
 - The Decoder parameters are also shared.
 - But its execution is task specific with each task token.
(this is similar to a task-specific expert chosen from a mixture of experts)
- Multi-task NLP-based transformers have natural language tokens as a task indicator (ex) summarization, translation,...), in the context of TAB, defined task tokens are used as indicators.

Structure

- Losses

- The model is trained with three losses
- \mathcal{L}_{clf} : classification loss (binary-cross entropy)
- \mathcal{L}_{kd} : knowledge distillation (which is applied on the probability)
 - Helps to reduce forgetting
- \mathcal{L}_{div} : divergence loss
 - Inspired from the auxiliary classifier of "DER"
 - Discriminate the current last task's classes \mathcal{C}^t and all previous classes
 - Encourages a better diversity among task tokens.
 - Discarded at test-time
- Total loss : $\mathcal{L} = (1 - \alpha)\mathcal{L}_{clf} + \alpha\mathcal{L}_{kd} + \lambda\mathcal{L}_{div}$
- α correspond to the fraction of the number of old classes over the number of new classes $\frac{|\mathcal{C}^{1:t-1}|}{|\mathcal{C}^{1:t}|}$


Experiments : CIFAR100, ImageNet100 and ImageNet1000

- Implementation details
 - 5 Self-Attention Blocks (SABs)
 - 1 Task-Attention Blocks (TAB)
 - Each Attention Block has 12 heads and 384 dimension embedding
- Hyperparameters are fixed for all experiments
← Use 10% of CIFAR 100's train set as validation and get hyperparameters.
(apply to all experiments equally)
- rehearsal memory : 2,000 images for CIFAR100 and ImageNet100, 20,000 images for ImageNet1000

Hyperparameter	CIFAR	ImageNet
# SAB		5
# CAB		1
# Attentions Heads		12
Embed Dim		384
Input Size	32	224
Patch Size	4	16

Table 1: **DyTox's architectures** for CIFAR and ImageNet. The only difference between the two architectures is the patch size, as the image sizes vary between datasets.

Experiments : Result of ImageNet100 and ImageNet1000



Methods	ImageNet100 10 steps					ImageNet1000 10 steps				
	#P	top-1		top-5		#P	top-1		top-5	
		Avg	Last	Avg	Last		Avg	Last	Avg	Last
ResNet18 joint	11.22	-	-	-	95.10	11.68	-	-	-	89.27
Transf. joint	11.00	-	79.12	-	93.48	11.35	-	73.58	-	90.60
<i>E2E</i> [5]	11.22	-	-	89.92	80.29	11.68	-	-	72.09	52.29
<i>Simple-DER</i> [48]	-	-	-	-	-	28.00	66.63	59.24	85.62	80.76
iCaRL [59]	11.22	-	-	83.60	63.80	11.68	38.40	22.70	63.70	44.00
BiC [32]	11.22	-	-	90.60	84.40	11.68	-	-	84.00	73.20
WA [81]	11.22	-	-	91.00	84.10	11.68	65.67	55.60	86.60	81.10
RPSNet [56]	-	-	-	87.90	74.00	-	-	-	-	-
DER w/o P [76]	112.27	77.18	66.70	93.23	87.52	116.89	68.84	60.16	88.17	82.86
DER [†] [76]	-	76.12	66.06	92.79	88.38	-	66.73	58.62	87.08	81.89
DyTox	11.01	77.15	69.10	92.04	87.98	11.36	71.29	63.34	88.59	84.49

The [†] marks the DER with setting specific pruning

Experiments : Result of CIFAR100

Methods	10 steps			20 steps			50 steps		
	#P	Avg	Last	#P	Avg	Last	#P	Avg	Last
ResNet18 Joint	11.22	-	80.41	11.22	-	81.49	11.22	-	81.74
Transf. Joint	10.72	-	76.12	10.72	-	76.12	10.72	-	76.12
iCaRL [59]	11.22	65.27 ± 1.02	50.74	11.22	61.20 ± 0.83	43.75	11.22	56.08 ± 0.83	36.62
UCIR [32]	11.22	58.66 ± 0.71	43.39	11.22	58.17 ± 0.30	40.63	11.22	56.86 ± 0.83	37.09
BiC [75]	11.22	68.80 ± 1.20	53.54	11.22	66.48 ± 0.32	47.02	11.22	62.09 ± 0.85	41.04
WA [81]	11.22	69.46 ± 0.29	53.78	11.22	67.33 ± 0.15	47.31	11.22	64.32 ± 0.28	42.14
PODNet [19]	11.22	58.03 ± 1.27	41.05	11.22	53.97 ± 0.85	35.02	11.22	51.19 ± 1.02	32.99
RPSNet [56]	56.5	68.60	57.05	-	-	-	-	-	-
DER w/o P [76]	112.27	75.36 ± 0.36	65.22	224.55	74.09 ± 0.33	62.48	561.39	72.41 ± 0.36	59.08
DER [†] [76]	-	74.64 ± 0.28	64.35	-	73.98 ± 0.36	62.55	-	72.05 ± 0.55	59.76
DyTox	10.73	73.66 ± 0.02	60.67 ± 0.34	10.74	72.27 ± 0.18	56.32 ± 0.61	10.77	70.20 ± 0.16	52.34 ± 0.26
DyTox+	10.73	75.54 ± 0.10	62.06 ± 0.25	10.74	75.04 ± 0.11	60.03 ± 0.45	10.77	74.35 ± 0.05	57.09 ± 0.13

DyTox constantly surpasses previous state-of-the-art model, despite having a comparable performance at the first step and fewer parameters.

Experiments

- DyTox is able to scale correctly while handling seamlessly the parameter growth by sharing most of the weights across tasks.
- In contrast, DER had to propose a complex pruning method, and this pruning required different hyperparameter values for different settings.
- Despite this, the pruning in DER[†] is less efficient when classes diversity increase: DER[†] doubles in size between ImageNet100 and ImageNet1000 while handling the same amount of tasks (10).

Improved training procedure

- To bridge the gap between DyTox and DER w/o P on CIFAR100, a method called "MixUp" is used to create new samples from existing samples
(DyTox+ : DyTox with MixUp)

- $\lambda \sim \text{Beta}(\alpha, \alpha)$ ($\alpha=0.8$)
 \rightarrow new images : $x = \lambda x_1 + (1 - \lambda)x_2$
as their labels : $y = \lambda y_1 + (1 - \lambda)y_2$

- Main Effect

- 1) It diversifies the training images and thus enlarges the training distribution in vicinity of each training sample
- 2) It improves the network calibration, reducing overconfidence in recent classes. (shared motivation with the sigmoid activation)

Training	1 step	50 steps	
	Last (\uparrow)	Last (\uparrow)	Forgetting (\downarrow)
DyTox	76.12	52.34	33.15
DyTox+	77.51 ^{+1.39}	57.09 ^{+4.75}	31.50 ^{-1.65}

Table 4: “**Last**” accuracy and forgetting [8] on CIFAR100 for the joint (1 step, no continual) and 50 steps settings.

Model Introspection : on CIFAR100

- Memory overhead
 - Only add a vector of size $d=384$ per task \rightarrow +0.004% per step
(setting of CIFAR100 with 50 tasks, memory overhead was almost null (+0.2%))
- Model ablations
 - Knowledge distillation
 - Fine-tuning :
 - Token Expansion
 - Divergence Classifier
 - Independent Classifier

		Knowledge Distillation	Finetuning	Token Expansion	Divergence Classifier	Indendepent Classifiers	Avg	Last
DyTox	Transformer						60.69	38.87
		✓					61.62	39.35
		✓	✓				63.42	42.21
	Dynamic	✓	✓	✓			67.30	47.57
		✓	✓	✓	✓		68.28	49.45
		✓	✓	✓	✓	✓	70.20	52.34

CIFAR100 with 50 steps

Conclusion

- Self-attention layers are shared across all tasks, and we add task-specific tokens to achieve task-specialized embeddings through a new task-attention layer.
- It has very little memory overhead and does not require complex hyperparameter tuning.
- The experiments show that the DyTox framework scales to large datasets like ImageNet1k with state-of-the-art performances and number of parameters of the model increases reasonably contrary to previous dynamic strategies.